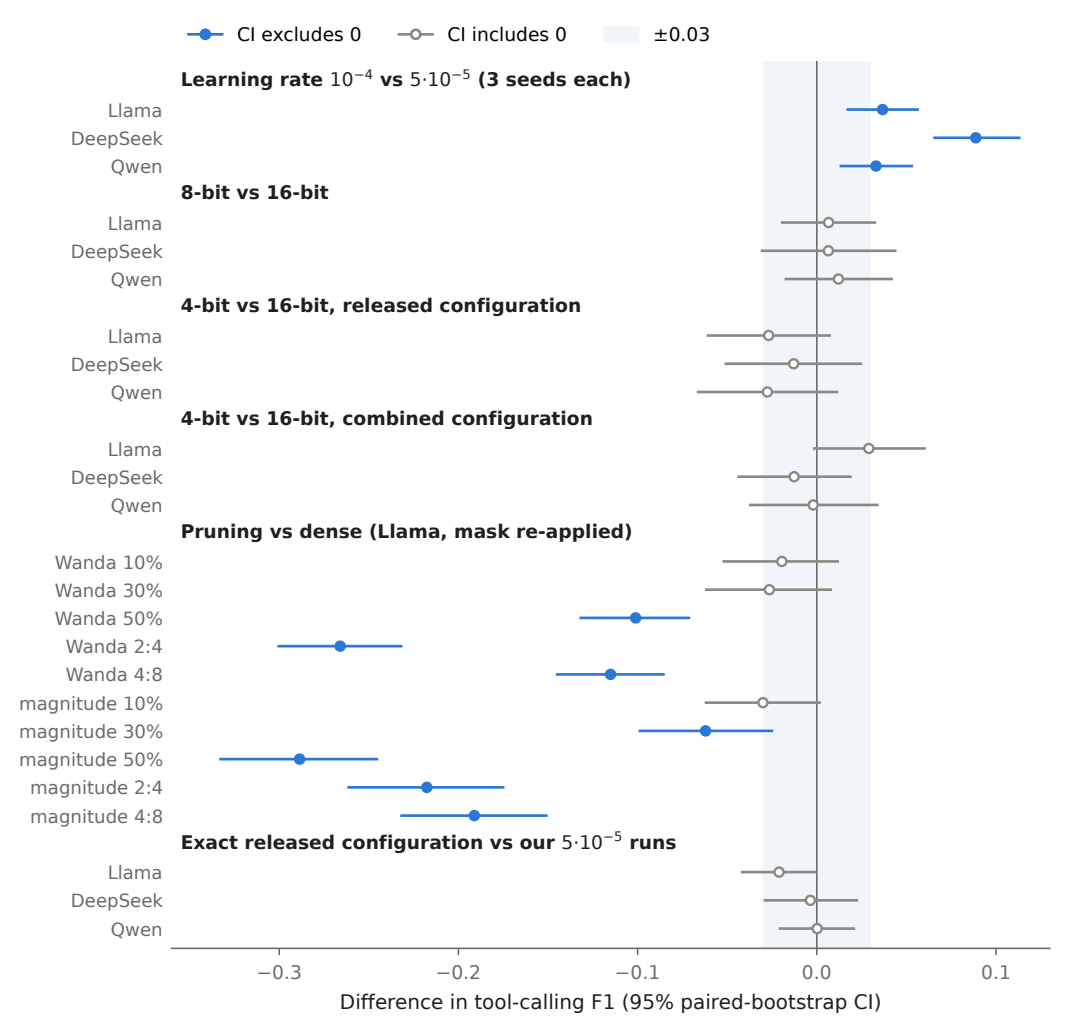
CI excludes 0
CI includes 0
±0.03
Learning rate 10^{-4} vs $5{\cdot}10^{-5}$ (3 seeds each)
Llama
DeepSeek
Qwen
8-bit vs 16-bit
Llama
DeepSeek
Qwen
4-bit vs 16-bit, released configuration
Llama
DeepSeek
Qwen
4-bit vs 16-bit, combined configuration
Llama
DeepSeek
Qwen
Pruning vs dense (Llama, mask re-applied)
Wanda 10%
Wanda 30%
Wanda 50%
Wanda 2:4
Wanda 4:8
magnitude 10%
magnitude 30%
magnitude 50%
magnitude 2:4
magnitude 4:8
Exact released configuration vs our $5{\cdot}10^{-5}$ runs
Llama
DeepSeek
Qwen
−0.3
−0.2
−0.1
0.0
0.1
Difference in tool-calling F1 (95% paired-bootstrap CI)