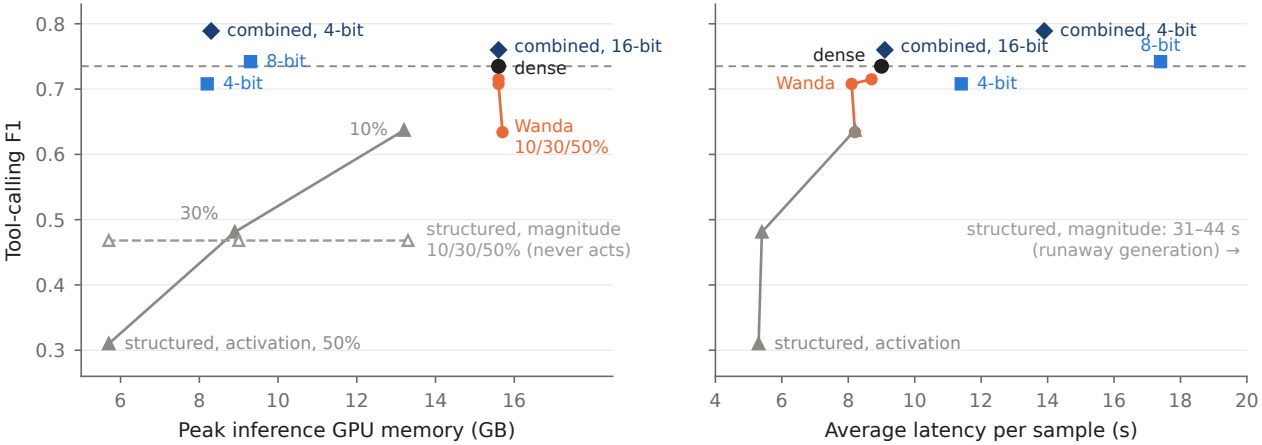
Tool-calling F1
0.8
0.7
0.6
0.5
0.4
0.3
combined, 4-bit
8-bit
4-bit
combined, 16-bit
dense
Wanda
10/30/50%
10%
30%
structured, magnitude
10/30/50% (never acts)
structured, activation, 50%
6
8
10
12
14
16
Peak inference GPU memory (GB)
combined, 4-bit
8-bit
dense
combined, 16-bit
Wanda
4-bit
structured, magnitude: 31–44 s
(runaway generation) →
structured, activation
4
6
8
10
12
14
16
18
20
Average latency per sample (s)